

Shiv Nadar University Chennai

Degree & Branch & Section	B.Tech Artificial Intelligence & Data Science - B	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27
Name & Reg. No.	Sadhumitha S.	24011101091

Experiment 1 Implementation of a Single Layer Perceptron for Binary Classification

1. Objective

The objective of this experiment is to understand the concept of an artificial neuron and implement a Single Layer Perceptron from scratch for binary classification. Students will learn the perceptron learning algorithm, understand the importance of activation functions, visualize the learning process, and evaluate the classifier using a real-world dataset.

2. Tasks to be Performed

1. Download and understand the dataset.
2. Perform Exploratory Data Analysis (EDA).
3. Preprocess the dataset.
4. Implement a Single Layer Perceptron from scratch.
5. Train the model using the perceptron learning rule.
6. Evaluate the classifier using standard performance metrics.
7. Analyze the effect of different learning rates.
8. Compare the implementation with Scikit-learn's Perceptron (optional).

3. Background Theory

Artificial Neuron is the basic building block of a neural network. It receives one or more input features, multiplies them with corresponding weights, adds a bias term, and applies an activation function to determine the output.

The Single Layer Perceptron is the first supervised learning algorithm that can solve linearly separable binary classification problems.

Mathematical Formulation

Weighted Sum

$$z = \mathbf{w}^T \mathbf{x} + b$$

where

- $\mathbf{x}$: Input vector
- $\mathbf{w}$: Weight vector
- b : Bias
- z : Linear combination

Prediction

$$\hat{y} = f(z)$$

Weight Update Rule

$$\mathbf{w}_{new} = \mathbf{w}_{old} + \eta(y - \hat{y})\mathbf{x}$$

Bias Update

$$b_{new} = b_{old} + \eta(y - \hat{y})$$

where η denotes the learning rate.

4. Activation Functions

Determines the output of a neuron based on the weighted sum of its inputs. It introduces non-linearity, enabling neural networks to learn complex patterns. Although the perceptron uses only the step activation function, modern deep learning models have differentiable activation functions for efficient optimization through backpropagation.

Step Activation Function

$$f(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$$

Produces a binary output and is suitable only for linearly separable binary classification problems.

Activation	Output Range	Typical Applications
Step	$\{0, 1\}$	Classical Perceptron
Sigmoid	$(0, 1)$	Binary Classification Output Layer
Tanh	$(-1, 1)$	Hidden Layers (Traditional Networks)
ReLU	$[0, \infty)$	Hidden Layers in CNNs and MLPs
Leaky ReLU	$(-\infty, \infty)$	Avoiding Dying ReLU
Softmax	$(0, 1)$	Multi-class Classification Output Layer

5. Dataset Description

Dataset: Banknote Authentication Dataset

Source: UCI Machine Learning Repository

Download Link:

<https://archive.ics.uci.edu/dataset/267/banknote+authentication>

Instances	1372
Features	4 Numerical Features
Classes	2
Missing Values	None
Task	Binary Classification

Feature Description

- Variance
- Skewness
- Curtosis
- Entropy

Target Class

- 0 – Authentic Banknote
- 1 – Forged Banknote

6. Experimental Procedure

Conducted using Python and the `slp.ipynb` Jupyter Notebook. The dataset was successfully loaded, preprocessed (standardized) and split into a 80-20 train-test ratio. The Single Layer Perceptron was implemented from scratch using NumPy and trained over several epochs using the perceptron learning rule. Visualizations for EDA and learning behaviour were generated and saved. Standard metrics like Accuracy, Precision, Recall and F1-score were also computed along with a Confusion Matrix.

7. Source Code

GitHub Repository

8. Experimental Results and Discussion

The perceptron successfully separated the classes and converged quickly due to the dataset being mostly linearly separable and features being normalized.

Performance Tables

Training Summary

Dataset Size	1372
Train/Test Split	1097 / 275
Learning Rate	0.01
Epochs	10 (Converged early if error = 0)
Accuracy	$\approx 97.8\%$
Precision	$\approx 98.4\%$
Recall	$\approx 96.85\%$
F1-score	$\approx 97.6\%$

Epoch-wise Learning (Sample)

Epoch	Errors	Weight 1	Weight 2	Bias
1	47	-0.0725	-0.0649	-0.0300
2	26	-0.0822	-0.0938	-0.0300
3	34	-0.0947	-0.0914	-0.0500

Generated Plots with Interpretation




Figure 1: Feature Histograms

Inference: The histograms reveal the underlying distributions of the four features. Notably, most features are not perfectly normally distributed and exhibit some skewness, suggesting that standardization is a crucial preprocessing step for the perceptron.




Figure 2: Correlation Heatmap

Inference: The correlation heatmap shows that the features have varying degrees of correlation with the Class. Variance and Skewness display significant negative correlation with the target class, making them highly informative features.




Figure 3: Scatter Plot of Variance vs Skewness

Inference: The scatter plot illustrates the class distribution in the feature space based on the two most highly correlated features. The classes appear to be mostly linearly separable with only a minor region of overlap, justifying the use of a Single Layer Perceptron.




Figure 4: Feature Boxplots

Inference: The boxplots highlight the spread and interquartile range of each feature. They

also show a few potential outliers in features like Curtosis and Entropy, though they are relatively minor.



Figure 5: Training Errors vs Epochs

Inference: The learning curve demonstrates a rapid decrease in the number of misclassifications over the epochs. The error reaches zero quickly, confirming that the perceptron successfully converged and found a separating hyperplane.




Figure 6: Confusion Matrix

Inference: The confusion matrix illustrates a highly accurate model with very few False Positives and False Negatives, demonstrating the strong classification capabilities of the trained perceptron on this dataset.




Figure 7: Learning Rate Comparison

Inference: This plot compares the convergence behavior across different learning rates (0.001, 0.01, 0.1).

It shows that while all models eventually converge, the rate of convergence and stability during the early epochs can vary depending on the chosen learning rate.

9. Additional Tasks

1. **Compare the Step and Sigmoid activation functions:** The Step function yields binary $\{0, 1\}$ outputs making it discontinuous and non-differentiable. Modern deep learning models rely on gradient descent and backpropagation, which necessitate differentiable activation functions like Sigmoid, Tanh, or ReLU.
2. **Compare with Scikit-learn's Perceptron:** Scikit-learn's Perceptron offers similar accuracy but provides optimized implementations, early stopping mechanisms, and regularization.
3. **Single Layer Perceptron and XOR:** The XOR problem is not linearly separable (no single straight line can separate the True and False outputs), hence a Single Layer Perceptron fails to solve it. It requires at least one hidden layer.
4. **Effect of Normalization:** Feature normalization ensures that all features contribute equally to the weight updates, thus leading to faster and more stable convergence.
5. **Implementation of Perceptron Learning Algorithm for Logic Gates:** The perceptron learning algorithm was implemented from scratch using Python for the OR, AND and NOT logic gates. The weights and bias were initialized to zero and updated according to the perceptron learning rule whenever a sample was misclassified. The updated weights and bias were displayed after every weight update during training. For the OR and AND gates, the decision boundary was plotted after each weight update to visualize the movement of the separating hyperplane. Since the NOT gate contains only one input feature, a 2D decision boundary is not applicable and hence, the learned weight and bias along with the final classification plot were presented.

OR Gate

The OR gate is linearly separable and therefore can be learned successfully by a Single Layer Perceptron. During training, the weights were updated whenever a misclassification occurred, causing the decision boundary to shift until all training samples were correctly classified.




Figure 8: OR Gate Decision Boundary after First Weight Update

Inference: Initially, the separating line does not correctly classify every sample. After the first weight update, the decision boundary begins moving toward the correct orientation.




Figure 9: Final Decision Boundary for OR Gate

Inference: The perceptron converges successfully and correctly classifies all four input combinations because the OR gate is linearly separable.

AND Gate

The AND gate is also linearly separable. Multiple weight updates gradually adjusted the separating hyperplane until only the input (1,1) was classified as the positive class.




Figure 10: AND Gate Decision Boundary after First Weight Update

Inference: The initial decision boundary incorrectly classifies several samples. The perceptron updates the weights to reduce these errors.




Figure 11: Final Decision Boundary for AND Gate

Inference: The final separating hyperplane correctly classifies all input combinations, demonstrating successful convergence of the perceptron.

NOT Gate

The NOT gate contains only one input feature. Consequently, a two-dimensional decision boundary cannot be plotted. Instead, the perceptron learns a single weight and bias that correctly invert the binary input.




Figure 12: Final Classification for NOT Gate

Inference: The perceptron correctly learns the NOT operation by assigning an appropriate weight and bias, producing the expected output for both input values.

Logic Gate Training Summary

Logic Gate	Linearly Separable	Perceptron Converged
OR	Yes	Yes
AND	Yes	Yes
NOT	Yes	Yes

10. Conclusion

In this experiment, a Single Layer Perceptron was successfully implemented from scratch for a binary classification task. The Banknote Authentication Dataset was explored, preprocessed, and used to train the perceptron. The model showed excellent performance with high accuracy and effectively demonstrated how the perceptron learning rule updates weights iteratively to find a linear decision boundary.

In the additional task, the perceptron learning algorithm was successfully demonstrated on the fundamental OR, AND and NOT logic gates. The weights after every update and the corresponding decision boundaries showed how the perceptron iteratively learns a separating hyperplane for linearly separable problems. This gave a clear understanding of the perceptron learning mechanism.

11. References

1. F. Rosenblatt, "The Perceptron," *Psychological Review*, 1958.
2. I. Goodfellow, Y. Bengio and A. Courville, *Deep Learning*, MIT Press, 2016.

3. C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006.
4. S. Haykin, *Neural Networks and Learning Machines*, Pearson, 2009.
5. UCI Machine Learning Repository – Banknote Authentication Dataset.
6. Scikit-learn Documentation: Perceptron.